

AI safety researcher building automated red-teaming, adversarial evals, and RL-based jailbreak discovery systems for agentic LLMs.

EDUCATION	Radboud University Nijmegen, Netherlands B.Sc. in Artificial Intelligence (Honours Programme) 2023 - 2026 (expected) Current GPA: 4.00/4.00 (Dutch system: 9/10).	
	Otto von Guericke University / Stendal University Germany B.A. Philosophy, Neuroscience & B.Sc. Psychology (Foundational Studies) 2020 - 2023 Consistently ranked in top percentile (GPA $\approx$ 3.9/4.0). Transferred to Radboud to specialize in technical AI.	
EXPERIENCE	Incoming Fellow LASR Labs Jul 2026 - Sep 2026 Selected for technical AI safety fellowship in London. Expected project focus: automated red-teaming, evals, and agentic LLM safety.	
	Independent Researcher & Foresight Grantee Foresight Institute Mar 2026 - Present - Awarded competitive \$22.5k AI Safety Grant to direct independent research on automated red-teaming and agentic LLM safety. - Developing an ICLR submission on split behavior in tool-using LLM agents, where models produce refusal-like text while still executing prohibited tool actions.	
	Student Researcher KachmanLab @ Radboud University Sep 2025 - Present - Developing novel jailbreaking techniques and adversarial attacks on LLMs. - Orchestrating multi-GPU distributed training (Slurm/HPC) for Clipped Importance Sampling Policy Optimization (CISPO) to test alignment boundaries using automated jailbreakers.	
	ARENA 5.0 Fellow Alignment Research Engineer Accelerator Apr 2025 - Jun 2025 - Selected for competitive fellowship with 4% acceptance rate, training alongside PhDs and industry professionals. - Implemented Transformers from scratch and conducted deep dives into RLHF and interpretability.	
	Research Assistant & Foresight Fellow Donders Institute / Foresight Institute Feb 2024 - Jul 2025 - Architected a scalable model-fitting pipeline on Slurm HPC clusters to quantify controllability, translating theoretical frameworks into executable code. - Selected as Foresight Fellow to lead computational modeling of agency, independently developing algorithms to extract latent parameters from behavioral data.	
	AI Safety Grantee & Scholar Long-Term Future Fund / AIMM Jul 2022 - Sep 2023 Awarded competitive grant (<5% acceptance) and mentorship to execute independent research under Jan Hendrik Kirchner (OpenAI/Anthropic).	
SELECTED RESEARCH	Split Behavior in Tool-Using LLM Agents 2026 ICLR submission in preparation. Studying cases where models produce refusal-like text while still executing prohibited or harmful tool actions. Building behavioral and mechanistic evidence that text-level refusal and action-level safety can decouple in agentic LLM systems.	
	David vs. Goliath: Verifiable Agent-to-Agent Jailbreaking via Reinforcement Learning 2026 NeurIPS 2026 submission / arXiv preprint. Led a project formalizing Tag-Along Attacks, where adversaries exploit tool necessity in agentic settings. Built an RL framework for automated jailbreak discovery, with attacks transferring zero-shot to closed models such as Gemini.	
	Reasoning Capabilities of LLMs 2025 EACL 2026 accepted paper from AI Safety Camp 2025. Contributed mechanistic interpretability analyses in an international research team studying internal representations of reasoning-related capabilities. Paper available on ACL Anthology.	

PROJECTS	Styx Interchange: Localizing Refusal Behavior in LLMs 2026
	<i>Weekend mechanistic interpretability sprint isolating refusal mechanisms in Qwen3-0.6B. Compared activation patching against input gradients using paired prompts, finding that causal refusal behavior is concentrated in specific model layers.</i>
	Internal Representations in SONAR Autoencoders 2025
	<i>ARENA 5.0 capstone. Analyzed internal representations in a text autoencoder to identify features correlated with code validity, grammaticality, and correctness-like structure.</i>
OPEN SOURCE	UK AISI inspect_ai: Merged PR fixing vLLM startup behaviour.
	Prime Intellect Environments: Merged PR for text retrieval environment. Implementing AI Safety benchmarks, e.g. AgentDojo environment.
SKILLS	Research: Automated red-teaming, adversarial evals, LLM jailbreaking, RL-based attack discovery, mechanistic interpretability.
	Engineering: PyTorch, Transformers, verifiers, Slurm/HPC, multi-GPU training, AgentDojo, inspect_ai.
LEADERSHIP & AWARDS	Other: Python, Pandas, MATLAB, Stan, JavaScript/React.
	Technical Advisor, Foresight AI Safety Grant Program2024
	User Experience Lead, Fridays for Future International2019-2020