

EDUCATION	Radboud University <i>B.Sc. in Artificial Intelligence (Honours Programme)</i> • Current GPA: 4.00/4.00 (Dutch system: 9/10). Top of class.	Nijmegen, Netherlands 2023 - 2026 (expected)
	Neuromatch Academy <i>Computational Neuroscience & Deep Learning</i> • Completed intensive 3-week curriculum on biologically plausible learning rules and high-dimensional data analysis.	Remote / Summer School Jul 2023
	Otto-Von-Guericke-University / Stendal University <i>B.A. Philosophy, Neuroscience & B.Sc. Psychology (Foundational Studies)</i> • Consistently ranked in top percentile (GPA $\approx$ 3.9/4.0). Transferred to Radboud to specialize in technical AI.	Germany 2020 - 2023
EXPERIENCE	Independent Researcher & Foresight Grantee Foresight Institute • Awarded competitive \$22.5k AI Safety Grant (<5% acceptance rate) to direct independent research for an upcoming NeurIPS 2026 submission. • Investigating the long-term behavior of latent adversarial training under continuous adversarial pressure.	Mar 2026 - Present
	Student Researcher KachmanLab @ Radboud University • Developing novel jailbreaking techniques and adversarial attacks on LLMs. • Orchestrating multi-GPU distributed training (Slurm/HPC) for Clipped Importance Sampling Policy Optimization (CISPO) to test alignment boundaries using automated jailbreakers.	Sep 2025 - Present
	ARENA 5.0 Fellow Alignment Research Engineer Accelerator • Selected for competitive fellowship with 4% acceptance rate, training alongside PhDs and industry professionals. • Implemented Transformers from scratch and conducted deep dives into RLHF and interpretability.	Apr 2025 - Jun 2025
	Research Assistant & Foresight Fellow Donders Institute / Foresight Institute • Architected a scalable model-fitting pipeline on Slurm HPC clusters to quantify controllability, translating theoretical frameworks into executable code. • Selected as Foresight Fellow to lead computational modeling of agency, independently developing algorithms to extract latent parameters from behavioral data.	Feb 2024 - Jul 2025
	AI Safety Grantee & Scholar Long-Term Future Fund / AIMM • Awarded competitive grant (<5% acceptance) and mentorship to execute independent research under Jan Hendrik Kirchner (OpenAI/Anthropic).	Jul 2022 - Sep 2023
PROJECTS	Styx Interchange: Localizing Refusal Behavior in LLMs <i>Executed a rapid weekend research sprint to isolate refusal mechanisms in Qwen3-0.6B. Compared activation patching against input gradients using paired prompts, demonstrating that causal refusal behavior is heavily concentrated in specific model layers.</i>	2026
	ICML2026 Submission: David vs. Goliath: Verifiable Agent-to-Agent Jailbreaking via Reinforcement Learning <i>A research project I led, formalizing Tag-Along Attacks, where adversaries exploit tool necessity. Our RL framework reveals that simple imperative syntax breaks safety guardrails (67% success) and transfers zero-shot to closed models like Gemini.</i>	2026
	ARENA 5.0 Capstone: Internal Representations in SONAR Autoencoders <i>Executed intensive 5-day research sprint in Mechanistic Interpretability. Analyzed model representations to identify features correlated with correctness.</i>	2025
	AI Safety Camp 2025: Reasoning Capabilities of LLMs <i>Mechanistic Interpretability research in an international team. Aiming for publication at NeurIPS/ICLM workshops. Paper available on ArXiv and accepted into EACL 2026.</i>	2025

OPEN SOURCE	UK AISI inspect_ai: Merged PR fixing vLLM startup behaviour.		
	Prime Intellect Environments: Merged PR for text retrieval environment. Implementing AI Safety benchmarks, e.g. AgentDojo environment.		
SKILLS	AI Safety & Research: LLM Jailbreaking, Adversarial Attacks (GRPO), Red Teaming, Representation Engineering, Mechanistic Interpretability.		
	Engineering: LLM Fine-Tuning (RLHF, Verifiers), Multi-GPU Distributed Training, Model Optimization.		
LEADERSHIP & AWARDS	Stack: Python (PyTorch, Pandas, Transformers, AgentDojo), Slurm/HPC, JavaScript (React), MATLAB, Stan, Java, Scala.		
	Technical Advisor, Foresight AI Safety Grant Program		2024
	User Experience Lead, Fridays for Future International		2019-2020
	Top of Year (Philosophy), A-levels		2019