

AI safety researcher building automated red-teaming, adversarial evals, and RL-based jailbreak discovery systems for agentic LLMs.

EDUCATION	Radboud University Nijmegen, Netherlands <i>B.Sc. in Artificial Intelligence (Honours Programme)</i> 2023 - 2026 (expected) Current GPA: 4.00/4.00 (Dutch system: 9/10).
	Otto von Guericke University / Stendal University Germany <i>B.A. Philosophy, Neuroscience & B.Sc. Psychology (Foundational Studies)</i> 2020 - 2023 Consistently ranked in top percentile (GPA $\approx$ 3.9/4.0). Transferred to Radboud to specialize in technical AI.
EXPERIENCE	Incoming Fellow LASR Labs Jul 2026 - Sep 2026 Selected for technical AI safety fellowship in London. Expected project focus: automated red-teaming, evals, and agentic LLM safety.
	Independent Researcher & Foresight Grantee Foresight Institute Mar 2026 - Present - Awarded competitive \$22.5k AI Safety Grant to direct independent research on automated red-teaming and agentic LLM safety. - Developing an ICLR submission on split behavior in tool-using LLM agents, where models produce refusal-like text while still executing prohibited tool actions.
	Student Researcher KachmanLab @ Radboud University Sep 2025 - Present - Developing novel jailbreaking techniques and adversarial attacks on LLMs. - Orchestrating multi-GPU distributed training (Slurm/HPC) for Clipped Importance Sampling Policy Optimization (CISPO) to test alignment boundaries using automated jailbreakers.
	ARENA 5.0 Fellow Alignment Research Engineer Accelerator Apr 2025 - Jun 2025 - Selected for competitive fellowship with 4% acceptance rate, training alongside PhDs and industry professionals. - Implemented Transformers from scratch and conducted deep dives into RLHF and interpretability.
	Research Assistant & Foresight Fellow Donders Institute / Foresight Institute Feb 2024 - Jul 2025 - Architected a scalable model-fitting pipeline on Slurm HPC clusters to quantify controllability, translating theoretical frameworks into executable code. - Selected as Foresight Fellow to lead computational modeling of agency, independently developing algorithms to extract latent parameters from behavioral data.
	AI Safety Grantee & Scholar Long-Term Future Fund / AIMM Jul 2022 - Sep 2023 Awarded competitive grant (<5% acceptance) and mentorship to execute independent research under Jan Hendrik Kirchner (OpenAI/Anthropic).
SELECTED RESEARCH	Split Behavior in Tool-Using LLM Agents 2026 <i>ICLR submission in preparation. Studying cases where models produce refusal-like text while still executing prohibited or harmful tool actions. Building behavioral and mechanistic evidence that text-level refusal and action-level safety can decouple in agentic LLM systems.</i>
	David vs. Goliath: Verifiable Agent-to-Agent Jailbreaking via Reinforcement Learning 2026 <i>NeurIPS 2026 submission / arXiv preprint. Led a project formalizing Tag-Along Attacks, where adversaries exploit tool necessity in agentic settings. Built an RL framework for automated jailbreak discovery, with attacks transferring zero-shot to closed models such as Gemini.</i>
	Reasoning Capabilities of LLMs 2025 <i>EACL 2026 accepted paper from AI Safety Camp 2025. Contributed mechanistic interpretability analyses in an international research team studying internal representations of reasoning-related capabilities. Paper available on ACL Anthology.</i>

PROJECTS	Styx Interchange: Localizing Refusal Behavior in LLMs 2026
	<i>Weekend mechanistic interpretability sprint isolating refusal mechanisms in Qwen3-0.6B. Compared activation patching against input gradients using paired prompts, finding that causal refusal behavior is concentrated in specific model layers.</i>
	Internal Representations in SONAR Autoencoders 2025
	<i>ARENA 5.0 capstone. Analyzed internal representations in a text autoencoder to identify features correlated with code validity, grammaticality, and correctness-like structure.</i>
OPEN SOURCE	UK AISI inspect_ai: Merged PR fixing vLLM startup behaviour.
	Prime Intellect Environments: Merged PR for text retrieval environment. Implementing AI Safety benchmarks, e.g. AgentDojo environment.
SKILLS	Research: Automated red-teaming, adversarial evals, LLM jailbreaking, RL-based attack discovery, mechanistic interpretability.
	Engineering: PyTorch, Transformers, verifiers, Slurm/HPC, multi-GPU training, AgentDojo, inspect_ai.
LEADERSHIP & AWARDS	Other: Python, Pandas, MATLAB, Stan, JavaScript/React.
	Technical Advisor, Foresight AI Safety Grant Program 2024
	User Experience Lead, Fridays for Future International 2019-2020